

## ▼ Imports

```
# Data Handling
import pandas as pd
import numpy as np

# Preprocessing & Splitting
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MinMaxScaler

# Machine Learning Models
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
import xgboost as xgb

# Model Optimisation (Fine-tuning Hyperparameters)
from sklearn.model_selection import (
    GridSearchCV,
    RandomizedSearchCV,
    StratifiedKFold
)

# Evaluation Metrics
from sklearn.metrics import (
    accuracy_score,
    precision_score,
    recall_score,
    f1_score,
    classification_report,
    precision_recall_fscore_support,
    confusion_matrix,
    ConfusionMatrixDisplay
)

# Visualisation Tools
import seaborn as sns
import matplotlib.pyplot as plt

# Model Explainer
import shap

# Save/Export Models and Datasets
import pickle
import joblib
from google.colab import files
```

## ✓ Preprocessing

## ✓ Loading the Dataset

```
# Load CSVs
```


```
features_df = pd.read_csv('datasets/elliptic_txs_features.csv')
```

```
classes_df = pd.read_csv('datasets/elliptic_txs_classes.csv')
```

```
↔ Downloading...  
From: https://drive.google.com/uc?id=1S04vme33DCGmL3VTYV9myph9suqLL71s  
To: /content/elliptic_txs_classes.csv  
100%|██████████| 3.31M/3.31M [00:00<00:00, 184MB/s]  
Downloading...  
From (original): https://drive.google.com/uc?id=1ZiHg7nmp00qt56xrp0XqB5zor6  
From (redirected): https://drive.google.com/uc?id=1ZiHg7nmp00qt56xrp0XqB5zo  
To: /content/elliptic_txs_features.csv  
100%|██████████| 690M/690M [00:06<00:00, 113MB/s]
```

```
# Quick preview
print("Features shape:", features_df.shape)
display(features_df.head())

print("Classes shape:", classes_df.shape)
display(classes_df.head())
```

 Features shape: (203768, 167)

	230425980	1	-0.1714692896288031	-0.18466755143291433	-1.2013688016765
0	5530458	1	-0.171484	-0.184668	-1.201
1	232022460	1	-0.172107	-0.184668	-1.201
2	232438397	1	0.163054	1.963790	-0.646
3	230460314	1	1.011523	-0.081127	-1.201
4	230459870	1	0.961040	-0.081127	-1.201

5 rows x 167 columns

Classes shape: (203769, 2)

	txId	class
0	230425980	unknown
1	5530458	unknown
2	232022460	unknown
3	232438397	2
4	230460314	unknown

## ✓ Cleaning the Dataset

```
# Assign proper names to transaction id and time step features
features_df.rename(columns={features_df.columns[0]: 'txId', features_df.columns

print("Features shape:", features_df.shape)
display(features_df.head())
```

 Features shape: (203768, 167)

	txId	timeStep	-0.1714692896288031	-0.18466755143291433	-1.2013688
0	5530458	1	-0.171484	-0.184668	
1	232022460	1	-0.172107	-0.184668	
2	232438397	1	0.163054	1.963790	
3	230460314	1	1.011523	-0.081127	
4	230459870	1	0.961040	-0.081127	

5 rows x 167 columns

```
# Merge the datasets
transactions_df = pd.merge(features_df, classes_df, on='txId', how='left')


print("Merged data shape:", transactions_df.shape)
transactions_df.head()
```

 Merged data shape: (203768, 168)

	txId	timeStep	-0.1714692896288031	-0.18466755143291433	-1.2013688
0	5530458	1	-0.171484	-0.184668	
1	232022460	1	-0.172107	-0.184668	
2	232438397	1	0.163054	1.963790	
3	230460314	1	1.011523	-0.081127	
4	230459870	1	0.961040	-0.081127	

5 rows x 168 columns


```
# Drop any rows with missing data
transactions_df = transactions_df.dropna()
display(transactions_df)
```



	txId	timeStep	-0.1714692896288031	-0.18466755143291433	-1.20
0	5530458	1	-0.171484	-0.184668	
1	232022460	1	-0.172107	-0.184668	
2	232438397	1	0.163054	1.963790	
3	230460314	1	1.011523	-0.081127	
4	230459870	1	0.961040	-0.081127	
...	...	...	...	...	...
203763	173077460	49	-0.145771	-0.163752	
203764	158577750	49	-0.165920	-0.123607	
203765	158375402	49	-0.172014	-0.078182	
203766	158654197	49	-0.172842	-0.176622	
203767	157597225	49	-0.012037	-0.132276	

203768 rows × 168 columns

```
# Find out all the values in class
transactions_df['class'].value_counts()
```



	count
class	
unknown	157204
2	42019
1	4545

dtype: int64

```
# Drop rows with unknown as the class
no_unknown_df = transactions_df.loc[(transactions_df['class'] != "unknown"), 't
transactions_df = transactions_df.loc[transactions_df['txId'].isin(no_unknown_c

display(transactions_df)
print("Transactions data shape:", transactions_df.shape)
```




	txId	timeStep	-0.1714692896288031	-0.18466755143291433	-1.20
2	232438397	1	0.163054	1.963790	
8	232029206	1	-0.005027	0.578941	
9	232344069	1	-0.147852	-0.184668	
10	27553029	1	-0.151357	-0.184668	
15	3881097	1	-0.172306	-0.184668	
...	...	...	...	...	...
203751	80329479	49	-0.159293	-0.037276	
203753	158406298	49	-0.172962	-0.126566	
203758	158375075	49	-0.170412	-0.078164	
203762	147478192	49	-0.093732	-0.116160	
203765	158375402	49	-0.172014	-0.078182	

46564 rows x 168 columns

Transactions data shape: (46564, 168)

```
# make class only 0s and 1s for binary classification
transactions_df['class'] = transactions_df['class'].replace('2', 0)
transactions_df['class'] = transactions_df['class'].replace('1', 1)
transactions_df['class'].value_counts()
```

 /tmp/ipython-input-9-1428541033.py:2: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: <https://pandas.pydata.org/pandas-docs>  
 transactions\_df['class'] = transactions\_df['class'].replace('2', 0)  
 /tmp/ipython-input-9-1428541033.py:3: FutureWarning: Downcasting behavior in  
 transactions\_df['class'] = transactions\_df['class'].replace('1', 1)  
 /tmp/ipython-input-9-1428541033.py:3: SettingWithCopyWarning:  
 A value is trying to be set on a copy of a slice from a DataFrame.  
 Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: <https://pandas.pydata.org/pandas-docs>  
 transactions\_df['class'] = transactions\_df['class'].replace('1', 1)


	count
class	
0	42019
1	4545

dtype: int64

## ✓ Splitting the Data into Train/Test Sets

```
# Split the features from the target and remove redundant features
X_raw = transactions_df.drop(columns=['txId', 'timeStep', 'class'])
y_raw = transactions_df['class']
```


```
print("X shape:", X_raw.shape)
print("y shape:", y_raw.shape)
```

 X shape: (46564, 165)  
y shape: (46564,)



```
# Split the data into 70% training, 30% testing
X_train_raw, X_test_raw, y_train, y_test = train_test_split(X_raw, y_raw, test_

display(X_train_raw.head())
print(f"Train set: {X_train_raw.shape}")
display(X_test_raw.head())
print(f"Test set: {X_test_raw.shape}")
```



	-0.1714692896288031	-0.18466755143291433	-1.2013688016765636	-0.1
147393	-0.147861	-0.016933	-0.091383	
33575	-0.159984	-0.184668	-1.201369	
103549	-0.167092	-0.158783	-1.201369	
130185	-0.172687	-0.024186	1.018602	
42840	-0.172408	-0.081127	-1.201369	

5 rows x 165 columns

Train set: (32594, 165)

	-0.1714692896288031	-0.18466755143291433	-1.2013688016765636	-0.1
42795	-0.172924	0.031421	0.463609	
182513	-0.058037	-0.113846	0.463609	
186795	-0.172945	-0.123312	1.018602	
32051	-0.156359	-0.184668	-1.201369	
86105	-0.172951	-0.175452	0.463609	

5 rows x 165 columns

Test set: (13970, 165)

## ✓ Scaling Numerical Features

```
# Scale the numerical features using MinMaxScaler
scaler = MinMaxScaler()
scaler.fit(X_train_raw)
X_train = scaler.transform(X_train_raw)
X_test = scaler.transform(X_test_raw)
```

## ✓ Initial Model Training and Evaluation

```
def evaluate_model(model, model_name="Model"):
    # Predict on train and test
    y_pred_train = model.predict(X_train)
    y_pred_test = model.predict(X_test)

    # Calculate metrics
    accuracy_train = accuracy_score(y_train, y_pred_train)
    accuracy_test = accuracy_score(y_test, y_pred_test)

    precision_train = precision_score(y_train, y_pred_train)
    precision_test = precision_score(y_test, y_pred_test)

    recall_train = recall_score(y_train, y_pred_train)
    recall_test = recall_score(y_test, y_pred_test)

    f1_train = f1_score(y_train, y_pred_train)
    f1_test = f1_score(y_test, y_pred_test)

    # Prepare results DataFrame
    results_df = pd.DataFrame([
        ['Accuracy', accuracy_train, accuracy_test],
        ['Precision', precision_train, precision_test],
        ['Recall', recall_train, recall_test],
        ['F1 Score', f1_train, f1_test]
    ], columns=['Metric', 'Train', 'Test'])

    print(f"{model_name} Results:")
    display(results_df)


    # Return metrics and predictions
    return results_df, y_pred_train, y_pred_test
```

## ✓ Logistic Regression

```
lr_model = LogisticRegression(max_iter=1000, random_state=0)
lr_model.fit(X_train, y_train)
```

```
results_lr, y_pred_train_lr, y_pred_test_lr = evaluate_model(lr_model, model_na
```

➞ Logistic Regression Results:


	Metric	Train	Test	
0	Accuracy	0.917408	0.917323	
1	Precision	0.689977	0.689655	
2	Recall	0.279157	0.278592	
3	F1 Score	0.397493	0.396867	

## ✓ Random Forest

```
rf_model = RandomForestClassifier(random_state=0)
rf_model.fit(X_train, y_train)
```

```
results_rf, y_pred_train_rf, y_pred_test_rf = evaluate_model(rf_model, model_na
```



➞ Random Forest Results:

	Metric	Train	Test	
0	Accuracy	1.0	0.987688	
1	Precision	1.0	0.994196	
2	Recall	1.0	0.879032	
3	F1 Score	1.0	0.933074	

## ✓ XGBClassifier

```
xgb_model = xgb.XGBClassifier(eval_metric='logloss', random_state=0)
xgb_model.fit(X_train, y_train)
results_xgb, y_pred_train_xgb, y_pred_test_xgb= evaluate_model(xgb_model, model
```

 XGBClassifier Results:

	Metric	Train	Test	
0	Accuracy	1.0	0.991410	
1	Precision	1.0	0.984424	
2	Recall	1.0	0.926686	
3	F1 Score	1.0	0.954683	

## ✓ Storing Results For Later Comparison

```
before_optimisation_results = {
    'Logistic Regression': results_lr,
    'Random Forest': results_rf,
    'XGBClassifier': results_xgb
}
```

## ✓ Model Optimisation (Fine-tuning Hyperparameters)

### ✓ Grid Search

```
def finetune_model_using_gridsearch(model, param_grid, X_train, y_train, scoring):
    if model_name:
        print(f"\nTuning hyperparameters for {model_name}...\n")

    grid_search = GridSearchCV(
        estimator=model,
        param_grid=param_grid,
        scoring=scoring,
        cv=StratifiedKFold(n_splits=5, shuffle=True, random_state=0),
        verbose=1,
        n_jobs=-1
    )

    grid_search.fit(X_train, y_train)

    print(f"\nBest parameters for {model_name}:\n{grid_search.best_params_}")
    print(f"Best {scoring} score: {grid_search.best_score_:.4f}\n")

    return grid_search.best_estimator_, grid_search
```

## ✓ Randomised Search

```
def finetune_model_using_randomisedsearch(model, param_grid, X_train, y_train,
    if model_name:
        print(f"\nTuning hyperparameters for {model_name}...\n")

    grid_search = RandomizedSearchCV(
        estimator=model,
        param_distributions=param_grid,
        n_iter=5,
        scoring=scoring,
        cv=StratifiedKFold(n_splits=4, shuffle=True, random_state=0),
        verbose=1,
        n_jobs=-1,
        random_state=0
    )

    grid_search.fit(X_train, y_train)


    print(f"\nBest parameters for {model_name}:\n{grid_search.best_params_}")
    print(f"Best {scoring} score: {grid_search.best_score_:.4f}\n")

    return grid_search.best_estimator_, grid_search
```

## ✓ Logistic Regression

```
# Defines parameter grid
param_grid_lr = {
    'tol': [1e-5, 1e-4, 1e-3, 1e-2],
    'solver': ['liblinear'],
    'C': [0.1, 1, 10],
    'penalty': ['l1', 'l2'],
    'max_iter': [50, 100, 300]
}

# Fine-tunes model
lr_best_model, lr_grid = finetune_model_using_randomisedsearch(
    LogisticRegression(max_iter=1000, random_state=0),
    param_grid_lr,
    X_train,
    y_train,
    model_name="Logistic Regression"
)
```

 Tuning hyperparameters for Logistic Regression...

Fitting 4 folds for each of 5 candidates, totalling 20 fits

Best parameters for Logistic Regression:  
 {'tol': 1e-05, 'solver': 'liblinear', 'penalty': 'l1', 'max\_iter': 50, 'C':  
 Best f1 score: 0.8039

/usr/local/lib/python3.11/dist-packages/sklearn/svm/\_base.py:1249: Converge  
 warnings.warn(

## ✓ Random Forest

```
# Defines parameter grid
param_grid_rf = {
    'n_estimators': [100, 200, 500],
    'max_depth': [10, 20, 50, None],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4],
    'max_features': ['sqrt', 'log2', None],
    'class_weight': ['balanced', None]
}

# Fine-tunes model
rf_best_model, rf_grid = finetune_model_using_randomisedsearch(
    model=RandomForestClassifier(random_state=0),
    param_grid=param_grid_rf,
    X_train=X_train,
    y_train=y_train,
    model_name="Random Forest"
)
```



Tuning hyperparameters for Random Forest...

Fitting 4 folds for each of 5 candidates, totalling 20 fits

Best parameters for Random Forest:

{'n\_estimators': 500, 'min\_samples\_split': 10, 'min\_samples\_leaf': 1, 'max\_

Best f1 score: 0.9306

## ✓ XGBClassifier

```
# Defines parameter grid
param_grid_xgb = {
    'n_estimators': [100, 300, 500],
    'max_depth': [3, 6, 10],
    'learning_rate': [0.01, 0.1, 0.2],
    'subsample': [0.6, 0.8, 1.0],
    'colsample_bytree': [0.6, 0.8, 1.0],
    'gamma': [0, 1, 5],
    'reg_alpha': [0, 0.1, 1],
    'reg_lambda': [0.1, 1, 10],
    'scale_pos_weight': [1, 10, 20]
}

# Fine-tunes model
xgb_best_model, xgb_grid = finetune_model_using_randomisedsearch(
    xgb.XGBClassifier(use_label_encoder=False, eval_metric='logloss', random_st
    param_grid_xgb,
    X_train,
    y_train,
    model_name="XGBClassifier"
)
```



Tuning hyperparameters for XGBClassifier...

Fitting 4 folds for each of 5 candidates, totalling 20 fits  
 /usr/local/lib/python3.11/dist-packages/xgboost/training.py:183: UserWarning  
 Parameters: { "use\_label\_encoder" } are not used.

```
bst.update(dtrain, iteration=i, fobj=obj)
```



Best parameters for XGBClassifier:  
 {'subsample': 1.0, 'scale\_pos\_weight': 20, 'reg\_lambda': 1, 'reg\_alpha': 0.  
 Best f1 score: 0.9500

## ✓ Re-training models using optimal parameters





```
results_lr, y_pred_train_lr, y_pred_test_lr = evaluate_model(lr_best_model, "Lc
```

 Logistic Regression (Optimised) Results:

	Metric	Train	Test	
0	Accuracy	0.965638	0.963064	
1	Precision	0.854245	0.830218	
2	Recall	0.781201	0.781525	
3	F1 Score	0.816092	0.805136	



```
results_rf, y_pred_train_rf, y_pred_test_rf = evaluate_model(rf_best_model, "Ra
```

 Random Forest (Optimised) Results:

	Metric	Train	Test	
0	Accuracy	0.996472	0.987115	
1	Precision	0.999023	0.983660	
2	Recall	0.964791	0.882698	
3	F1 Score	0.981609	0.930448	

```
results_xgb, y_pred_train_xgb, y_pred_test_xgb = evaluate_model(xgb_best_model,
```

 XGBClassifier (Optimised) Results:

	Metric	Train	Test	
0	Accuracy	0.999540	0.991696	
1	Precision	0.995307	0.960857	
2	Recall	1.000000	0.953812	
3	F1 Score	0.997648	0.957322	

```
# Storing Results For Later Comparison
after_optimisation_results = {
    'Logistic Regression': results_lr,
    'Random Forest': results_rf,
    'XGBClassifier': results_xgb
}
```

## ✓ Performance Results and Comparison

### ✓ Performance Before vs After Model Optimisation

```
def compile_split_comparison_tables(before_results, after_results, model_name):
    metrics = ['Accuracy', 'Precision', 'Recall', 'F1 Score']

    train_before = before_results[model_name]['Train'].values
    train_after = after_results[model_name]['Train'].values
    train_delta = train_after - train_before

    test_before = before_results[model_name]['Test'].values
    test_after = after_results[model_name]['Test'].values
    test_delta = test_after - test_before

    train_df = pd.DataFrame({
        'Metric': metrics,
        'Train (Before)': train_before,
        'Train (After)': train_after,
        'Train Δ': train_delta
    })

    test_df = pd.DataFrame({
        'Metric': metrics,
        'Test (Before)': test_before,
        'Test (After)': test_after,
        'Test Δ': test_delta
    })

    return train_df, test_df


for model_name in before_optimisation_results.keys():
    print(f"\nTrain Comparison for {model_name}")
    train_table, test_table = compile_split_comparison_tables(before_optimisation_results[model_name],
                                                              after_optimisation_results[model_name])
    display(train_table)

    print(f"\nTest Comparison for {model_name}")
    display(test_table)
```





Train Comparison for Logistic Regression



	Metric	Train (Before)	Train (After)	Train Δ	
0	Accuracy	0.917408	0.965638	0.048230	

1	Precision	0.689977	0.854245	0.164269	
2	Recall	0.279157	0.781201	0.502043	
3	F1 Score	0.397493	0.816092	0.418599	



## Test Comparison for Logistic Regression

	Metric	Test (Before)	Test (After)	Test $\Delta$	
0	Accuracy	0.917323	0.963064	0.045741	
1	Precision	0.689655	0.830218	0.140563	
2	Recall	0.278592	0.781525	0.502933	
3	F1 Score	0.396867	0.805136	0.408269	



## Train Comparison for Random Forest

	Metric	Train (Before)	Train (After)	Train $\Delta$	
0	Accuracy	1.0	0.996472	-0.003528	
1	Precision	1.0	0.999023	-0.000977	
2	Recall	1.0	0.964791	-0.035209	
3	F1 Score	1.0	0.981609	-0.018391	

## Test Comparison for Random Forest


	Metric	Test (Before)	Test (After)	Test $\Delta$	
0	Accuracy	0.987688	0.987115	-0.000573	
1	Precision	0.994196	0.983660	-0.010536	
2	Recall	0.879032	0.882698	0.003666	
3	F1 Score	0.933074	0.930448	-0.002626	

## Train Comparison for XGBClassifier

	Metric	Train (Before)	Train (After)	Train $\Delta$	
0	Accuracy	1.0	0.999540	-0.000460	
1	Precision	1.0	0.995307	-0.004693	
2	Recall	1.0	1.000000	0.000000	
3	F1 Score	1.0	0.997648	-0.002352	

## Test Comparison for XGBClassifier

	Metric	Test (Before)	Test (After)	Test $\Delta$	
--	--------	---------------	--------------	---------------	---

0	Accuracy	0.991410	0.991696	0.000286	
1	Precision	0.984424	0.960857	-0.023567	
2	Recall	0.926686	0.953812	0.027126	
3	F1 Score	0.954683	0.957322	0.002639	

Next  
steps:

[Generate code with train\\_table](#)
[View recommended plots](#)
[New interactive sheet](#)

```
def plot_model_before_after_bars(model_name, before_results, after_results):
    metrics = ['Accuracy', 'Precision', 'Recall', 'F1 Score']

    # Extract values
    train_before = before_results[model_name]['Train'].values
    train_after = after_results[model_name]['Train'].values
    test_before = before_results[model_name]['Test'].values
    test_after = after_results[model_name]['Test'].values

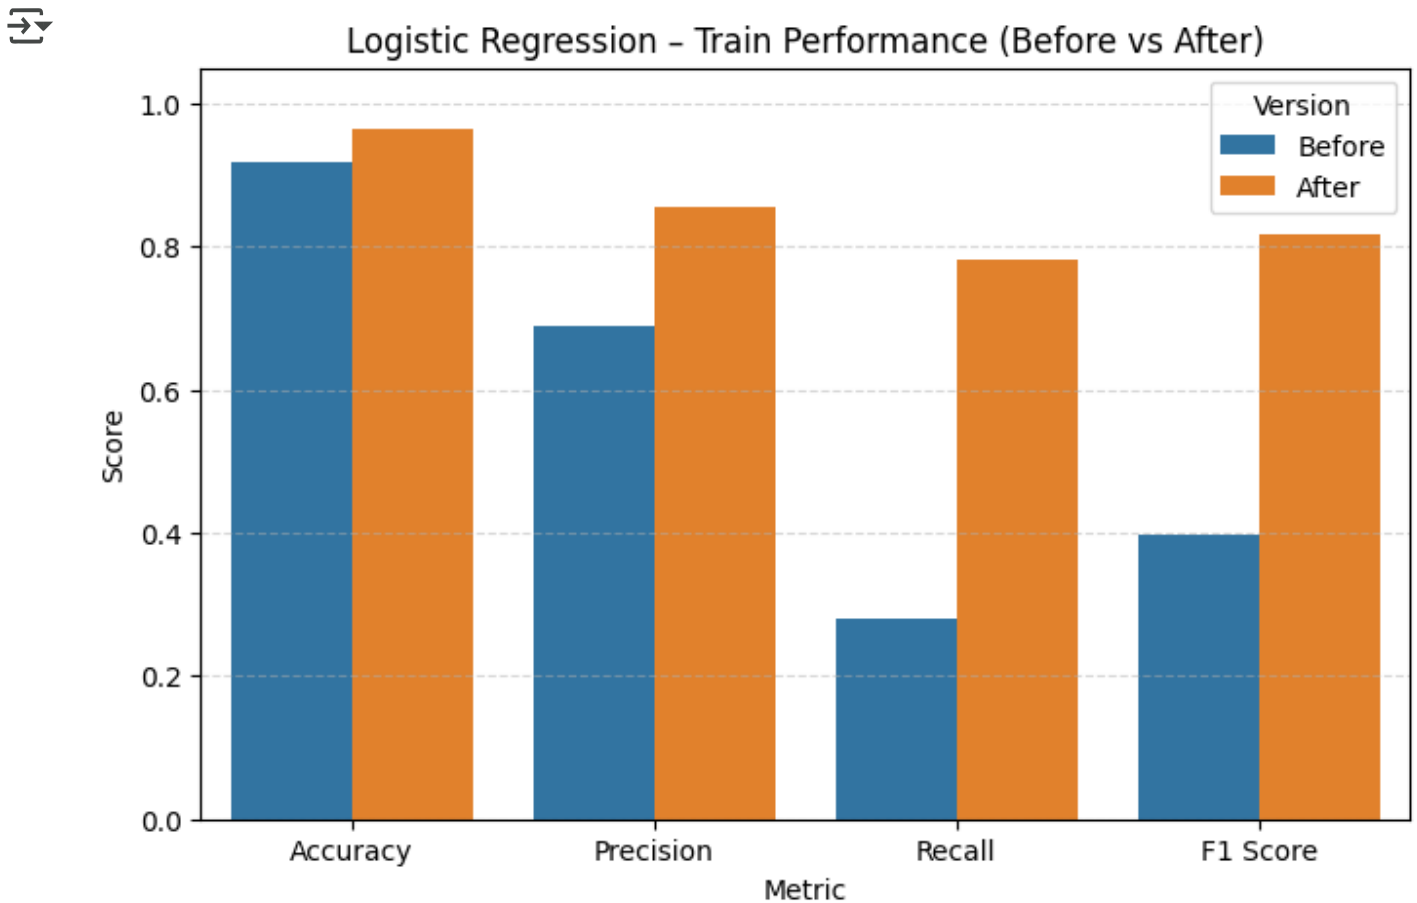
    # Prepare data for plotting
    train_data = pd.DataFrame({
        'Metric': metrics,
        'Before': train_before,
        'After': train_after
    }).melt(id_vars='Metric', var_name='Version', value_name='Score')

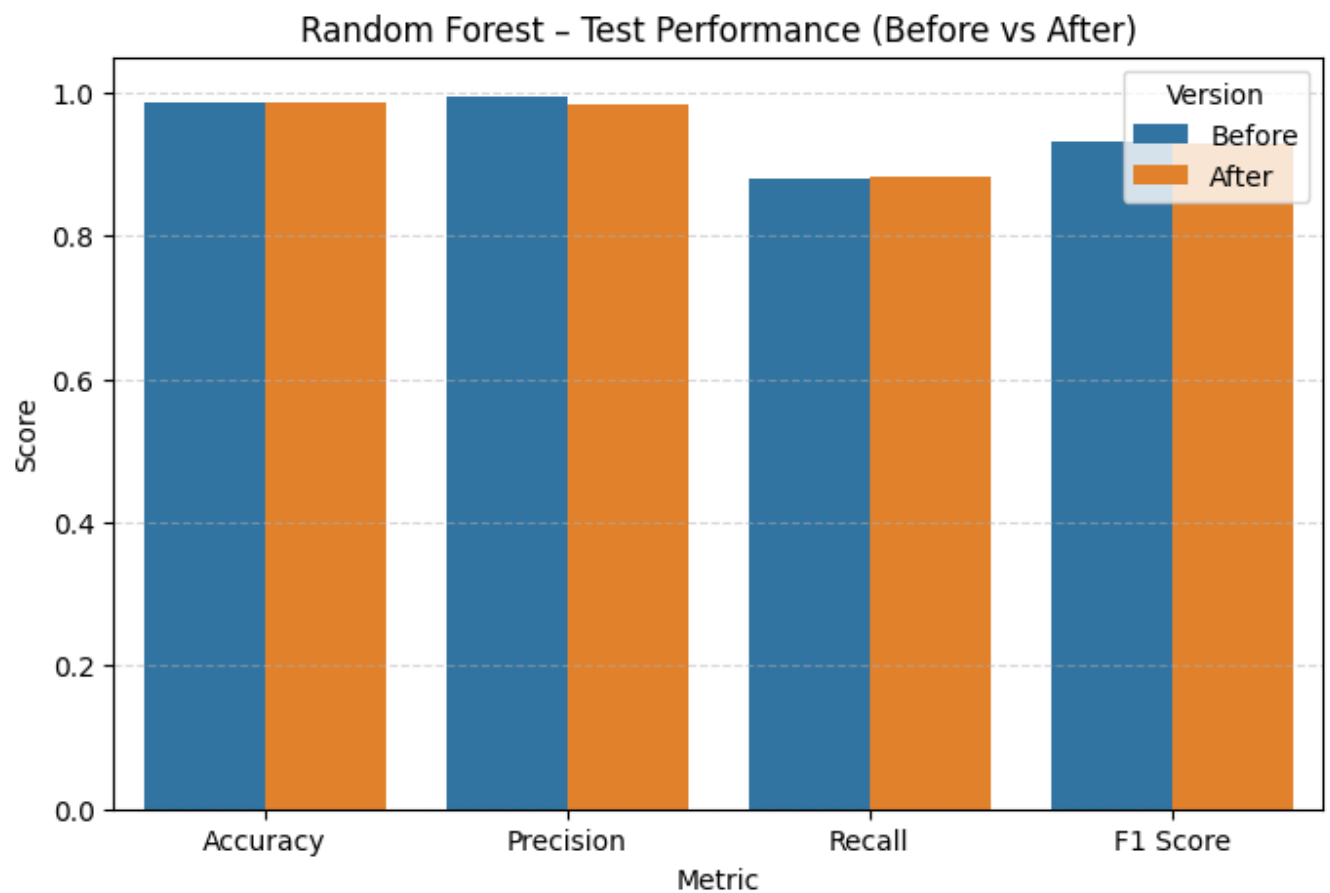

    test_data = pd.DataFrame({
        'Metric': metrics,
        'Before': test_before,
        'After': test_after
    }).melt(id_vars='Metric', var_name='Version', value_name='Score')

    # Plot Train
    plt.figure(figsize=(8, 5))
    sns.barplot(data=train_data, x='Metric', y='Score', hue='Version')
    plt.title(f'{model_name} - Train Performance (Before vs After)')
    plt.ylim(0, 1.05)
    plt.grid(axis='y', linestyle='--', alpha=0.5)
    plt.show()

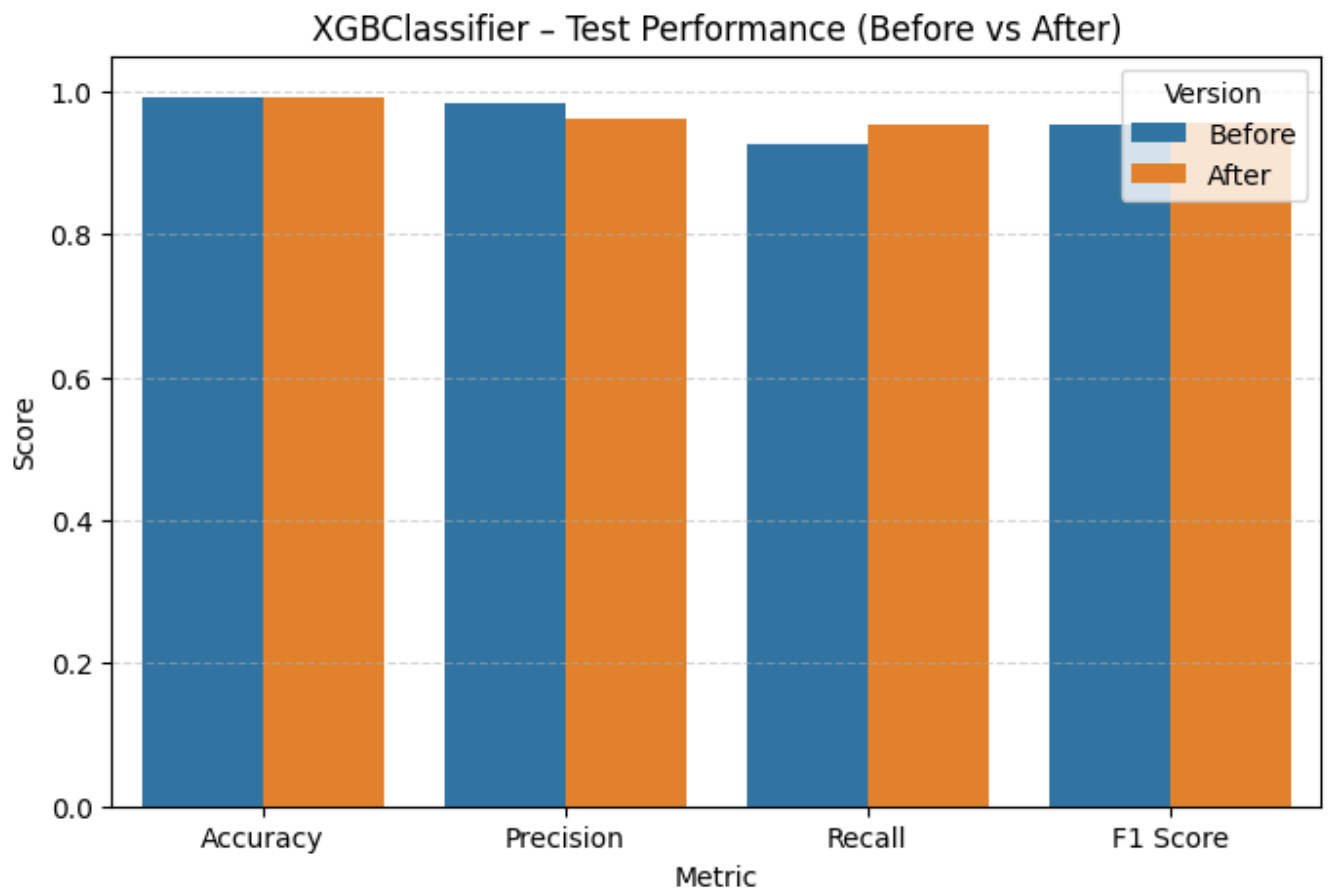

    # Plot Test
    plt.figure(figsize=(8, 5))
    sns.barplot(data=test_data, x='Metric', y='Score', hue='Version')
    plt.title(f'{model_name} - Test Performance (Before vs After)')
    plt.ylim(0, 1.05)
    plt.grid(axis='y', linestyle='--', alpha=0.5)
    plt.show()
```

```
for model_name in before_optimisation_results.keys():  
    plot_model_before_after_bars(model_name, before_optimisation_results, after
```





XGBClassifier - Train Performance (Before vs After)



## ✓ Comparing Performance Between Models

```
def create_metric_table(metric_name, model_results_dict):
    rows = []

    for model_name, df in model_results_dict.items():
        # Extract train and test values for the given metric
        metric_row = df[df['Metric'] == metric_name]
        if not metric_row.empty:
            train_val = metric_row['Train'].values[0]
            test_val = metric_row['Test'].values[0]
            rows.append([model_name, train_val, test_val])

    # Create DataFrame and sort by Test performance
    df_metric = pd.DataFrame(rows, columns=['Model', 'Train', 'Test'])
    df_metric_sorted = df_metric.sort_values(by='Test', ascending=False).reset_


    return df_metric_sorted

# Show sorted comparison tables for each metric
for metric in ['Accuracy', 'Precision', 'Recall', 'F1 Score']:
    print(f"\n{metric} Comparison")
    display(create_metric_table(metric, after_optimisation_results))
```





## Accuracy Comparison

	Model	Train	Test	
0	XGBClassifier	0.999540	0.991696	
1	Random Forest	0.996472	0.987115	
2	Logistic Regression	0.965638	0.963064	

## Precision Comparison

	Model	Train	Test	
0	Random Forest	0.999023	0.983660	
1	XGBClassifier	0.995307	0.960857	
2	Logistic Regression	0.854245	0.830218	

## Recall Comparison

	Model	Train	Test	
0	XGBClassifier	1.000000	0.953812	
1	Random Forest	0.964791	0.882698	
2	Logistic Regression	0.781201	0.781525	

## F1 Score Comparison

	Model	Train	Test	
0	XGBClassifier	0.997648	0.957322	
1	Random Forest	0.981609	0.930448	
2	Logistic Regression	0.816092	0.805136	

```
def plot_metric_bar_chart(metric_name, model_results_dict):
```

```
    data = []
```

```
    for model_name, df in model_results_dict.items():
```

```
        row = df[df['Metric'] == metric_name]
```

```
        if not row.empty:
```

```
            train_val = row['Train'].values[0]
```

```
            test_val = row['Test'].values[0]
```

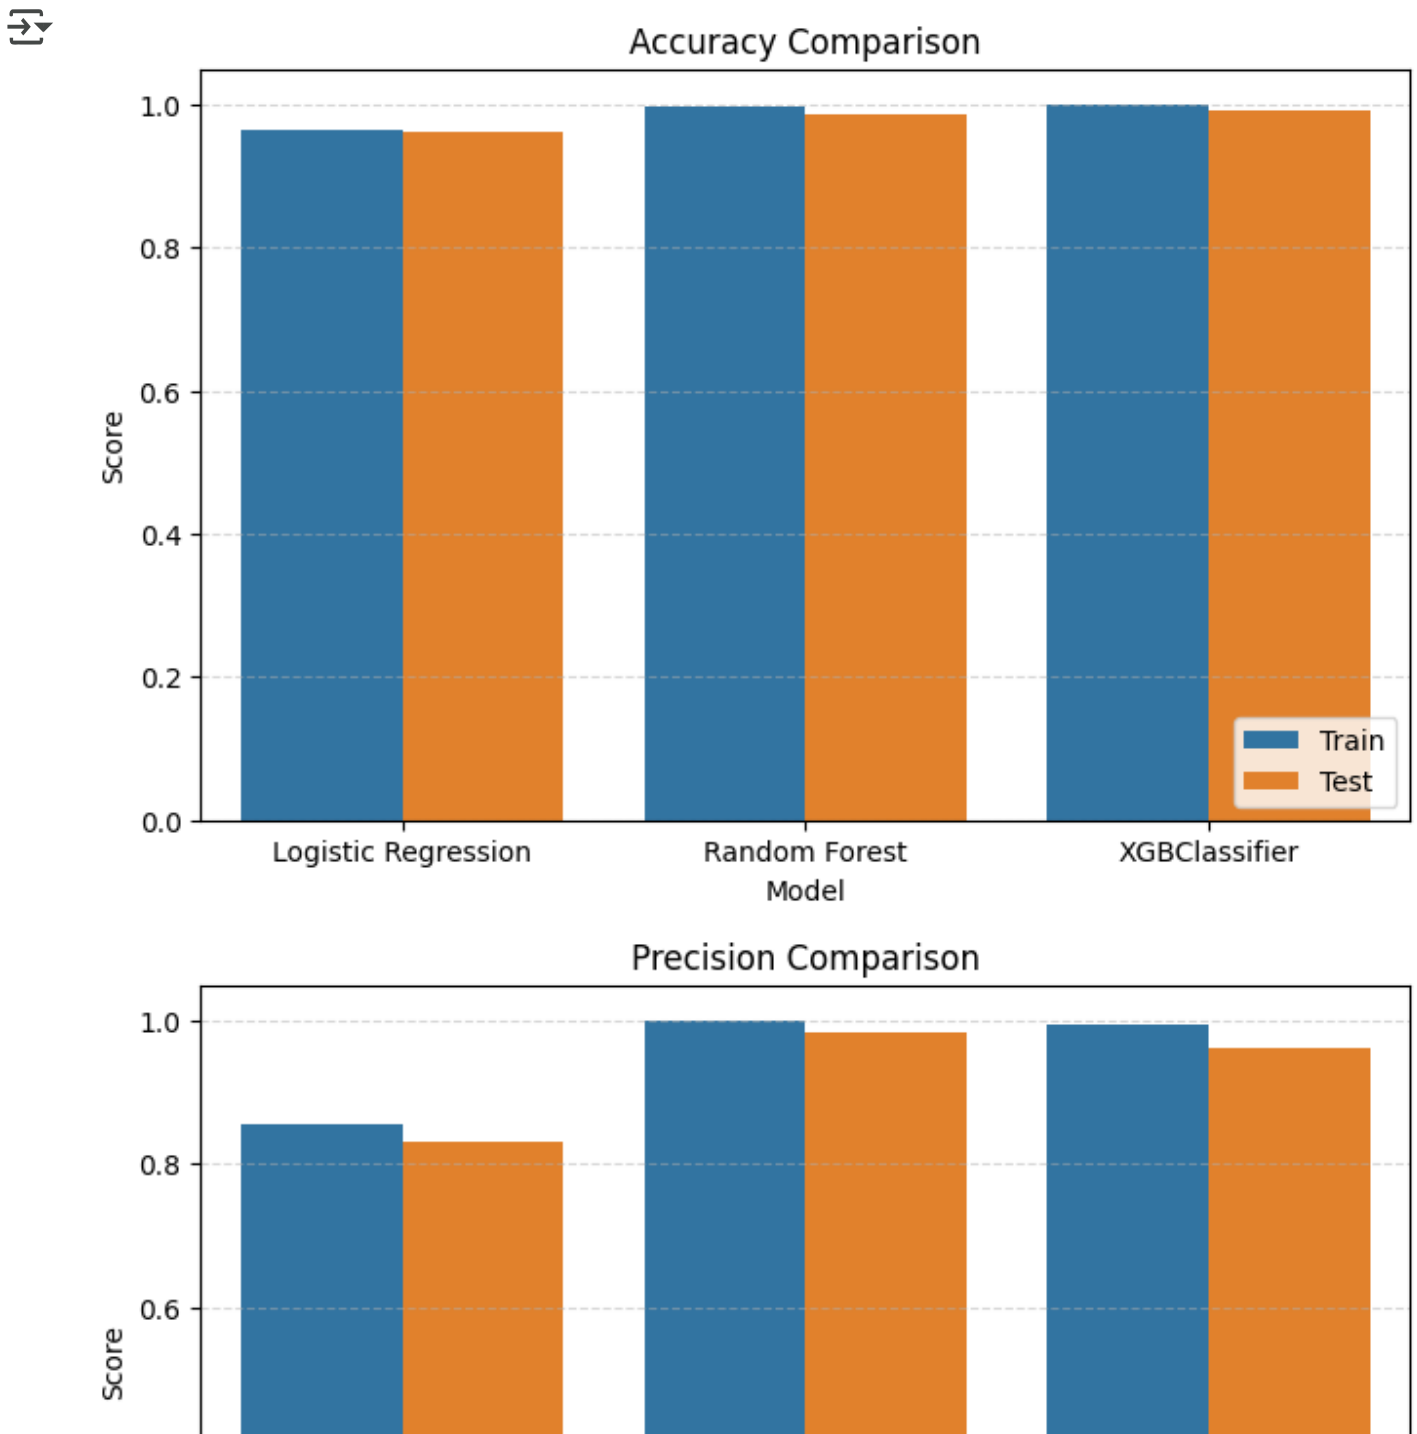
```
            data.append({'Model': model_name, 'Set': 'Train', 'Score': train_val})
```




```
            data.append({'Model': model_name, 'Set': 'Test', 'Score': test_val})
```


```
df_plot = pd.DataFrame(data)

plt.figure(figsize=(8, 5))
sns.barplot(data=df_plot, x='Model', y='Score', hue='Set')
plt.title(f'{metric_name} Comparison')
plt.ylim(0, 1.05)
plt.legend(loc='lower right')
plt.grid(axis='y', linestyle='--', alpha=0.5)
plt.show()

for metric in ['Accuracy', 'Precision', 'Recall', 'F1 Score']:
    plot_metric_bar_chart(metric, after_optimisation_results)
```








## ✓ Confusion Matrix For Each Model


```
def plot_confusion_matrix(y_true, y_pred, model_name):
    disp = ConfusionMatrixDisplay.from_predictions(
        y_true, y_pred,
        display_labels=["Class 0", "Class 1"],
        cmap='Blues',
        normalize=None
    )
    disp.ax_.set_title(f'{model_name} - Confusion Matrix')
    plt.show()

plot_confusion_matrix(y_train, y_pred_train_lr, "Logistic Regression (Train)")
plot_confusion_matrix(y_test, y_pred_test_lr, "Logistic Regression (Test)")


plot_confusion_matrix(y_train, y_pred_train_rf, "Random Forest (Train)")
plot_confusion_matrix(y_test, y_pred_test_rf, "Random Forest (Test)")

plot_confusion_matrix(y_train, y_pred_train_xgb, "XGBClassifier (Train)")
plot_confusion_matrix(y_test, y_pred_test_xgb, "XGBClassifier (Test)")
```






Logistic Regression (Test) - Confusion Matrix




Random Forest (Train) - Confusion Matrix







Random Forest (Test) - Confusion Matrix



XGBClassifier (Train) - Confusion Matrix





- ✓ Export for Deployment
- ✓ Model to be Deployed

```
# Save the best model
joblib.dump(xgb_best_model, 'model.pkl')
```

```
# Download/export the best model
files.download('model.pkl')
```

 Downloading "model.pkl": 

## ✓ SHAP Explainer

```
# Create SHAP explainer for your XGBoost model
explainer = shap.TreeExplainer(xgb_model)
```

```
# Save explainer to file
with open('explainer.pkl', 'wb') as f:
    pickle.dump(explainer, f)
```

```
# Download/export the SHAP explainer
files.download('explainer.pkl')
```

 Downloading "explainer.pkl": 